Regularization

K. Breininger, V. Christlein, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle, M. Zinnen, K. Packhäuser

Pattern Recognition Lab, Friedrich-Alexander University of Erlangen-Nürnberg

January 4, 2023

Tasks in this exercise

1. Optimization Constraints: Augmenting the loss function

2. Dropout **Layer**

3. Batch Normalization **Layer**

4. LeNet: Put everything together **(optional)**

5. RNN layer: Elman Unit

6. LSTM layer: Backpropagation at its best! **(optional)**

Optimization Constraints: Loss function augmentation

General outline

- Constraints change the total loss ...
- ... and have influence on the weight update of the respective layer!

General outline

- Constraints change the total loss ...

- ... and have influence on the weight update of the respective layer!

- Implement constraints as separate classes

→ **Independent** of loss function

General outline

- Constraints change the total loss ...

- ... and have influence on the weight update of the respective layer!

- Implement constraints as separate classes
- → **Independent** of loss function

- Constraints **only need** current weights
- → Add constraint objects in the optimizer

General outline

- Constraints change the total loss ...

- ... and have influence on the weight update of the respective layer!

- Implement constraints as separate classes
→ **Independent** of loss function

- Constraints **only need** current weights
→ Add constraint objects in the optimizer

- Since constraints generate part of the loss:
→ Change Neural Network container class (and associated classes) to "channel" and gather **regularization loss** for **all layers**

Workflow

- Forward pass

→ Calculate norm of weights in each trainable layer and gather as regularization loss

→ Add regularization loss to the final loss

Workflow

- Forward pass

→ Calculate norm of weights in each trainable layer and gather as regularization loss

→ Add regularization loss to the final loss

- Backward pass

→ In each trainable layer, include **the gradient of norm** when calculating update

L_2 regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta\lambda)\mathbf{w}^{(k)}}_{\text{Shrinkage}} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

L_2 regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda\|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta\lambda)\mathbf{w}^{(k)}}_{\text{Shrinkage}} - \eta\frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

- In the Forward pass the L2 norm gets squared, which eliminates the square root inside and increases the numerical stability as the gradient is easier to compute.

L_2 regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta\lambda)\mathbf{w}^{(k)}}_{\text{Shrinkage}} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

- In the Forward pass the L2 norm gets squared, which eliminates the square root inside and increases the numerical stability as the gradient is easier to compute.

- Notice for matrices we compute here the Frobenius norm, not the Spectral norm.

L$_2$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda\|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta\lambda)\mathbf{w}^{(k)}}_{\text{Shrinkage}} - \eta\frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

- In the Forward pass the L2 norm gets squared, which eliminates the square root inside and increases the numerical stability as the gradient is easier to compute.

- Notice for matrices we compute here the Frobenius norm, not the Spectral norm.

- The influence of constraints is controlled via λ. Because `lambda` is a python keyword, you want to use e.g. `alpha` instead.

L$_1$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda\|\mathbf{w}\|_1$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{\mathbf{w}^{(k)} - \eta\lambda\mathrm{sign}\left(\mathbf{w}^{(k)}\right)}_{\text{Other shrinkage}} - \eta\frac{\partial L}{\partial\mathbf{w}^{(k)}}$$
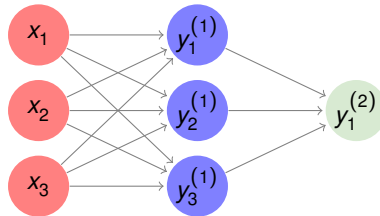
Dropout

Method

Figure: Dropout

- Implement this as a **fixed-function layer**

Method


Figure: Dropout

- Implement this as a **fixed-function layer**
- Randomly set **activations** $\mapsto 0$ with probability $1 - p$

Method

Figure: Dropout

- Implement this as a **fixed-function layer**
- Randomly set **activations** $\mapsto 0$ with probability $1 - p$
- **Test-time**: multiply activations with p

Inverted Dropout

- Can we get rid of the dropout layer at test-time?

Inverted Dropout

- Can we get rid of the dropout layer at test-time?
- → Change the behavior during training
- Multiply activations in forward-pass **only during training** by $\frac{1}{p}$
- Note: the backward pass has to be adapted as well!

Batch normalization

Forward pass

→ Normalization as a new layer with 2 parameters, γ and β

Forward pass

→ Normalization as a new layer with 2 parameters, γ and β

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \boldsymbol{\mu}_B}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}}$$

$\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ from **batch**

Forward pass

→ Normalization as a new layer with 2 parameters, $\boldsymbol{\gamma}$ and $\boldsymbol{\beta}$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \boldsymbol{\mu}_B}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}}$$

$\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ from **batch**

$$\hat{\mathbf{Y}} = \boldsymbol{\gamma}\tilde{\mathbf{X}} + \boldsymbol{\beta}$$

Forward pass

→ Normalization as a new layer with 2 parameters, $\boldsymbol{\gamma}$ and $\boldsymbol{\beta}$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \boldsymbol{\mu}_B}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}}$$

$\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ from **batch**

$$\hat{\mathbf{Y}} = \boldsymbol{\gamma}\tilde{\mathbf{X}} + \boldsymbol{\beta}$$

- $\boldsymbol{\mu}$, $\boldsymbol{\sigma}^2$ have the **same dimension** as the **input vectors**

Forward pass

→ Normalization as a new layer with 2 parameters, $\boldsymbol{\gamma}$ and $\boldsymbol{\beta}$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \boldsymbol{\mu}_B}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}}$$

$\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ from **batch**

$$\hat{\mathbf{Y}} = \boldsymbol{\gamma}\tilde{\mathbf{X}} + \boldsymbol{\beta}$$

- $\boldsymbol{\mu}$, $\boldsymbol{\sigma}^2$ have the **same dimension** as the **input vectors**
- $\boldsymbol{\beta}$, $\boldsymbol{\gamma}$ and $\boldsymbol{\mu}_B$, $\boldsymbol{\sigma}_B^2$ have same **dimension** to be able to preserve **identity**

Forward pass

→ Normalization as a new layer with 2 parameters, $\boldsymbol{\gamma}$ and $\boldsymbol{\beta}$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \boldsymbol{\mu}_B}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}}$$

$\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ from **batch**

$$\hat{\mathbf{Y}} = \boldsymbol{\gamma}\tilde{\mathbf{X}} + \boldsymbol{\beta}$$

- $\boldsymbol{\mu}$, $\boldsymbol{\sigma}^2$ have the **same dimension** as the **input vectors**
- $\boldsymbol{\beta}$, $\boldsymbol{\gamma}$ and $\boldsymbol{\mu}_B$, $\boldsymbol{\sigma}_B^2$ have same **dimension** to be able to preserve **identity**
- Notice that $\boldsymbol{\beta}$ is a **bias**

Test time

- Test-time: replace $\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ with $\boldsymbol{\mu}$ and $\boldsymbol{\sigma}^2$ of the **training set**

Test time

- Test-time: replace μ_B and σ_B^2 with μ and σ^2 of the **training set**
- It's **expensive** to calculate the true training set mean and variance

Test time

- Test-time: replace $\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ with $\boldsymbol{\mu}$ and $\boldsymbol{\sigma}^2$ of the **training set**
- It's **expensive** to calculate the true training set mean and variance
- Therefore a **moving average** is common:

$$\tilde{\boldsymbol{\mu}}^{(k)} \approx \alpha\tilde{\boldsymbol{\mu}}^{(k-1)} + (1-\alpha)\boldsymbol{\mu}_B^{(k)}$$
$$\tilde{\boldsymbol{\sigma}}^{\mathbf{2}(k)} \approx \alpha\tilde{\boldsymbol{\sigma}}^{\mathbf{2}(k-1)} + (1-\alpha)\boldsymbol{\sigma}_B^{\mathbf{2}(k)}$$

Test time

- Test-time: replace $\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B^2$ with $\boldsymbol{\mu}$ and $\boldsymbol{\sigma}^2$ of the **training set**
- It's **expensive** to calculate the true training set mean and variance
- Therefore a **moving average** is common:

$$\tilde{\boldsymbol{\mu}}^{(k)} \approx \alpha\tilde{\boldsymbol{\mu}}^{(k-1)} + (1-\alpha)\boldsymbol{\mu}_B^{(k)}$$
$$\tilde{\boldsymbol{\sigma}}^{2(k)} \approx \alpha\tilde{\boldsymbol{\sigma}}^{2(k-1)} + (1-\alpha)\boldsymbol{\sigma}_B^{2(k)}$$

- Moving average **decay** α (e.g. 0.8)
- The exponent (k) and (k-1) are iteration-indices!

Backward pass

- Gradient **with respect to weights** is simply:

$$\frac{\partial L}{\partial \gamma} = \sum_{b=1}^{B} \frac{\partial L}{\partial \hat{\mathbf{Y}}_b} \tilde{\mathbf{X}}_b = \sum_{b=1}^{B} \mathbf{E}_b \tilde{\mathbf{X}}_b$$

- For the **bias** likewise we have:

$$\frac{\partial L}{\partial \beta} = \sum_{b=1}^{B} \frac{\partial L}{\partial \hat{\mathbf{Y}}_b} = \sum_{b=1}^{B} \mathbf{E}_b$$

Backward pass

The **gradient with respect to the input** is more complicated, but here it is:

$$\frac{\partial L}{\partial \tilde{\mathbf{X}}} = \frac{\partial L}{\partial \hat{\mathbf{Y}}} \odot \boldsymbol{\gamma}$$

$$\frac{\partial L}{\partial \boldsymbol{\sigma}_B^2} = \sum_{b=1}^{B} \frac{\partial L}{\partial \tilde{\mathbf{X}}_b} \odot (\mathbf{X}_b - \boldsymbol{\mu}_B) \odot \frac{-1}{2} \left(\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon} \right)^{\frac{-3}{2}}$$

$$\frac{\partial L}{\partial \boldsymbol{\mu}_B} = \left(\sum_{b=1}^{B} \frac{\partial L}{\partial \tilde{\mathbf{X}}_b} \odot \frac{-1}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}} \right) + \underbrace{\frac{\partial L}{\partial \boldsymbol{\sigma}_B^2} \odot \frac{\sum_{b=1}^{B} -2(\mathbf{X}_b - \boldsymbol{\mu}_B)}{B}}_{0}$$

$$\frac{\partial L}{\partial \mathbf{X}} = \frac{\partial L}{\partial \tilde{\mathbf{X}}} \odot \frac{1}{\sqrt{\boldsymbol{\sigma}_B^2 + \boldsymbol{\epsilon}}} + \frac{\partial L}{\partial \boldsymbol{\sigma}_B^2} \odot \frac{2(\mathbf{X} - \boldsymbol{\mu}_B)}{B} + \frac{\partial L}{\partial \boldsymbol{\mu}_B} \odot \frac{1}{B}$$

Backward pass

- $\odot$ denotes an element-wise multiplication. Always check the dimensionality of your matrices!

Backward pass

- $\odot$ denotes an element-wise multiplication. Always check the dimensionality of your matrices!

- To make life easier, we will provide the code for the computation of the gradient with respect to the input:

Backward pass

- $\odot$ denotes an element-wise multiplication. Always check the dimensionality of your matrices!

- To make life easier, we will provide the code for the computation of the gradient with respect to the input:

- `compute_bn_gradients`

Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**

Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar μ, σ is calculated for the H **channels**

Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**

- A scalar μ, σ is calculated for the H **channels**

- Implementation can be reused, by observing
  - → that **spatial dimensions** M, N can be treated **like the batch dimension** B

Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**

- A scalar μ, σ is calculated for the H **channels**

- Implementation can be reused, by observing
  - → that **spatial dimensions** M, N can be treated **like the batch dimension** B
  - → we can **reshape** the $B \times H \times M \times N$ tensor to $B \times H \times M \cdot N$
  - → because of our format we have to **transpose** from $B \times H \times M \cdot N$ to $B \times M \cdot N \times H$
  - → and afterwards **reshape again** to have a $B \cdot M \cdot N \times H$ tensor

Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar μ, σ is calculated for the H **channels**
- Implementation can be reused, by observing
  - → that **spatial dimensions** M, N can be treated **like the batch dimension** B
  - → we can **reshape** the $B \times H \times M \times N$ tensor to $B \times H \times M \cdot N$
  - → because of our format we have to **transpose** from $B \times H \times M \cdot N$ to $B \times M \cdot N \times H$
  - → and afterwards **reshape again** to have a $B \cdot M \cdot N \times H$ tensor
- Consequently we have to **reverse this** before returning the **output**

Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar μ, σ is calculated for the H **channels**
- Implementation can be reused, by observing
  - → that **spatial dimensions** M, N can be treated **like the batch dimension** B
  - → we can **reshape** the $B \times H \times M \times N$ tensor to $B \times H \times M \cdot N$
  - → because of our format we have to **transpose** from $B \times H \times M \cdot N$ to $B \times M \cdot N \times H$
  - → and afterwards **reshape again** to have a $B \cdot M \cdot N \times H$ tensor
- Consequently we have to **reverse this** before returning the **output**
- ... and do the **same** in the **backward pass**
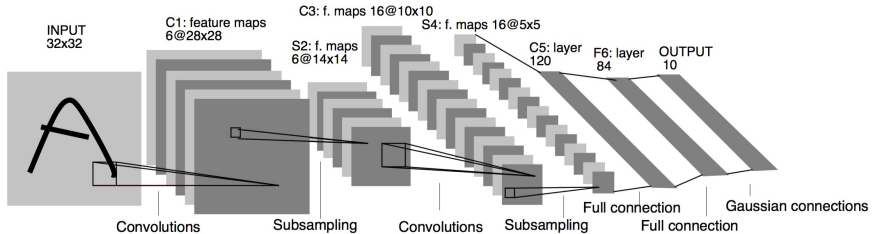
LeNet (optional)

LeNet architecture

Figure: LeNet

Modified LeNet architecture

Deviations

- Input is 28 × 28
- Our conv only supports "same" padding - so C3 has **larger activation maps**
- Input to **C5** is also **larger**
- We only implemented ReLUs, so **no** TanH
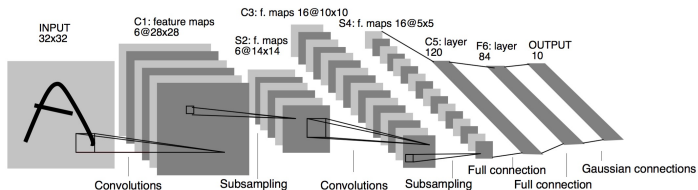- We also use the implemented SoftMax **instead of** RBF units

Figure: LeNet

Thanks for listening.
Any questions?